

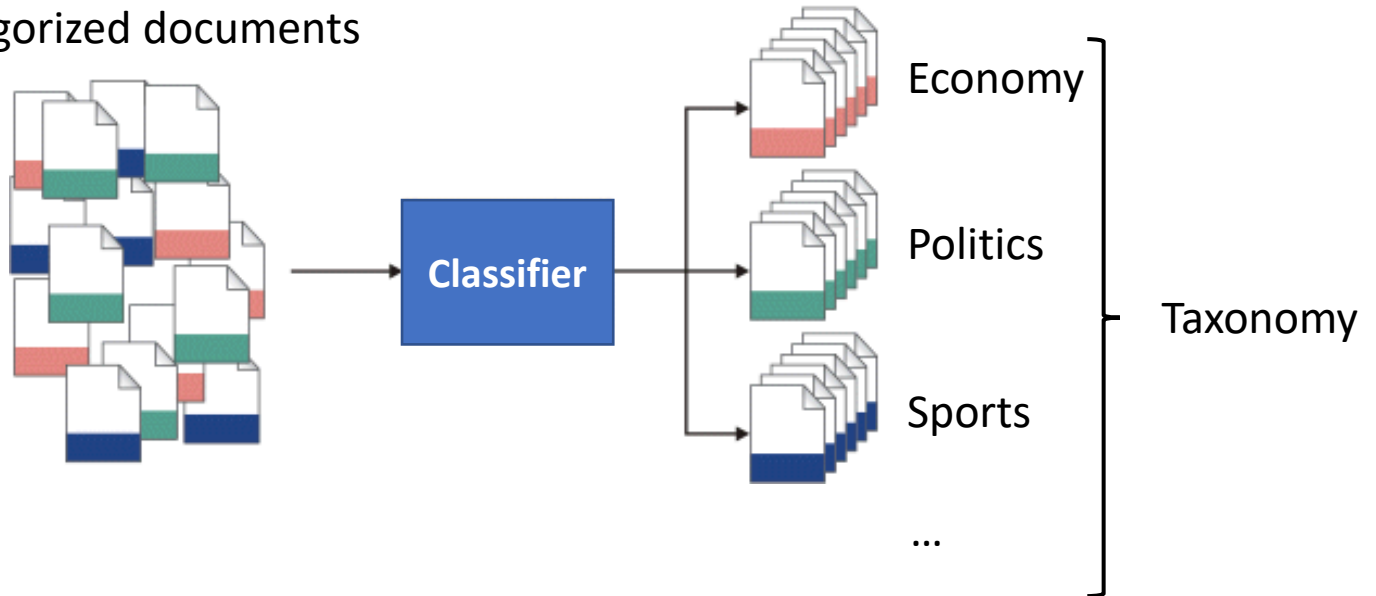
Document Categorization

Perceptron, Topic Detection, Sentiment Classification

Information Retrieval

Document Topic Categorization

Uncategorized documents



Spam filtering: Another text classification task

From: "" <takworld@hotmail.com>

Subject: real estate is the only way... gem oalvgkay

Anyone can buy real estate with no money down

Stop paying rent TODAY !

There is no need to spend hundreds or even thousands for similar courses

I am 22 years old and I have already purchased 6 properties using the methods outlined in this truly INCREDIBLE ebook.

Change your life NOW !

=====

Click Below to order:

<http://www.wholesaledaily.com/sales/nmd.htm>

=====

Sentiment Classification



THE AMAZING SPIDER-MAN

PG-13 , 2h 16m
[adventure](#) , [mystery and thriller](#) , [action](#)
Directed By: [Marc Webb](#)
In Theaters: [Jul 3, 2012](#) [Wide](#)
Streaming: [Sep 1, 2014](#)
[Marvel Studios](#)



The Amazing Spider-Man: Trailer - Four Minute Super Preview

4 minutes 2 seconds
Added: Feb 8, 2017



The Amazing Spider-Man: Official Clip - Those Are the Best Kind

1 minute 47 seconds
Added: Nov 5, 2019



The Amazing Spider-Man: Official Clip - High School Attack

2 minutes 59 seconds
Added: Nov 5, 2019

[VIEW ALL VIDEOS \(10\)](#)

THE AMAZING SPIDER-MAN REVIEWS

[All Critics](#)

[Top Critics](#)

[All Audience](#)

[NEXT →](#)



Daniel K

★★★★★

Sep 19, 2020

For years I had the blu ray sitting in my collection. I had seen amazing spider man 2 multiple times. I can confirm this movie is amazing. The whole time I was watching I wasn't bored. The movie is entertaining and the score is phenomenal. The first 45 minutes is slow but once uncle Ben dies (spoiler alert) it picks right up. Andrew Garfield is Spider-Man. At first he uses his powers to go after uncle Bens killer but that plot line is abandon and I didn't mind it. After he saves a little boy in a car wreck he takes his...

[Show More](#)



Jacob B

★★★★☆

Sep 14, 2020

As a whole, I like it. Some of the high school scenes are a little too goofy for me and quite honestly don't feel like Marc Webb directed them at all. It also feels as if some important content was missing from the movie (which has been confirmed). Other than that, I like the movie.



France Carl C

★★★★★

Sep 10, 2020

Andrew Garfield is the best Spider-Man



Samantha S

★★★★☆

Sep 04, 2020

Also confused for apparently not having rated this.



E C

★★★★☆

Sep 02, 2020

I thought this film was decent, but it was definitely not as good as the Sam Raimi "Spider-Man" films, and I also didn't think it was as good as the ones with Tom Holland. I thought this film was way too dark in some scenes in my opinion.



Target Sentiment on Twitter

- [Twitter Sentiment App](#)
- Alec Go, Richa Bhayani, Lei Huang. 2009. Twitter Sentiment Classification using Distant Supervision

Type in a word and we'll highlight the good and the bad

The figure displays sentiment analysis results for the target "united airlines". It consists of two charts: a pie chart titled "Sentiment by Percent" and a horizontal bar chart titled "Sentiment by Count".

Sentiment by Percent: A pie chart showing the distribution of sentiment. The red slice represents Negative sentiment at 68%, and the green slice represents Positive sentiment at 32%.

Sentiment by Count: A horizontal bar chart showing the count of tweets for each sentiment. The green bar represents Positive sentiment with a count of 11, and the red bar represents Negative sentiment with a count of 23.

Sentiment	Count	Percent
Positive	11	32%
Negative	23	68%

jljacobson: OMG... Could @United airlines have worse customer service? W8g now 15 minutes on hold 4 questions about a flight 2DAY that need a human.
Posted 2 hours ago

12345clumsy6789: I hate United Airlines Ceiling!!! Fukn impossible to get my conduit in this damn mess! ?
Posted 2 hours ago

EMLandPRGbelgiu: EML/PRG fly with Q8 united airlines and 24seven to an exotic destination. <http://t.co/Z9QloAjF>
Posted 2 hours ago

CountAdam: FANTASTIC customer service from United Airlines at XNA today. Is tweet more, but cell phones off now!
Posted 4 hours ago

Importance of information categorization

P. Jackson and I. Moulinier. 2002. *Natural Language Processing for Online Applications*

- “There is no question concerning the commercial value of being able to classify documents automatically by content. There are myriad potential applications of such a capability for corporate intranets, government departments, and Internet publishers”
- “Understanding the data is one of the keys to successful categorization, yet this is an area in which most categorization tool vendors are extremely weak. Many of the ‘one size fits all’ tools on the market have not been tested on a wide range of content types.”

Categorization/Classification

- Given:
 - A description of an instance, $d \in X$
 - X is the *instance language* or *instance space*.
 - Issue: how to represent text documents.
 - Usually some type of high-dimensional space
 - A fixed set of classes:
 $C = \{c_1, c_2, \dots, c_J\}$
- Determine:
 - The category of d : $\gamma(d) \in C$, where $\gamma(d)$ is a *classification function* whose domain is X and whose range is C .
 - We want to know how to build classification functions (“classifiers”).

Supervised Classification

- Given:
 - A description of an instance, $d \in X$
 - X is the *instance language* or *instance space*.
 - A fixed set of classes:
 $C = \{c_1, c_2, \dots, c_J\}$
 - A training set D of labeled documents with each labeled document $\langle d, c \rangle \in X \times C$
- Determine:
 - A learning method or algorithm which will enable us to learn a classifier $\gamma: X \rightarrow C$
 - For a test document d , we assign it the class $\gamma(d) \in C$

Domain specific taxonomies

- Domain specific terminologies are curated by domain experts and are designed with specific tasks and workflows in mind.
- In the medical domain, the SNOMED-CT is intended to describe medical conditions, procedures, admin, etc.
 - <http://browser.ihtsdotools.org/>
- In the computer science domain the ACM Computing Classification Scheme is widely used to classify published articles.
 - <https://dl.acm.org/ccs/ccs.cfm>



Wikipedia as a database

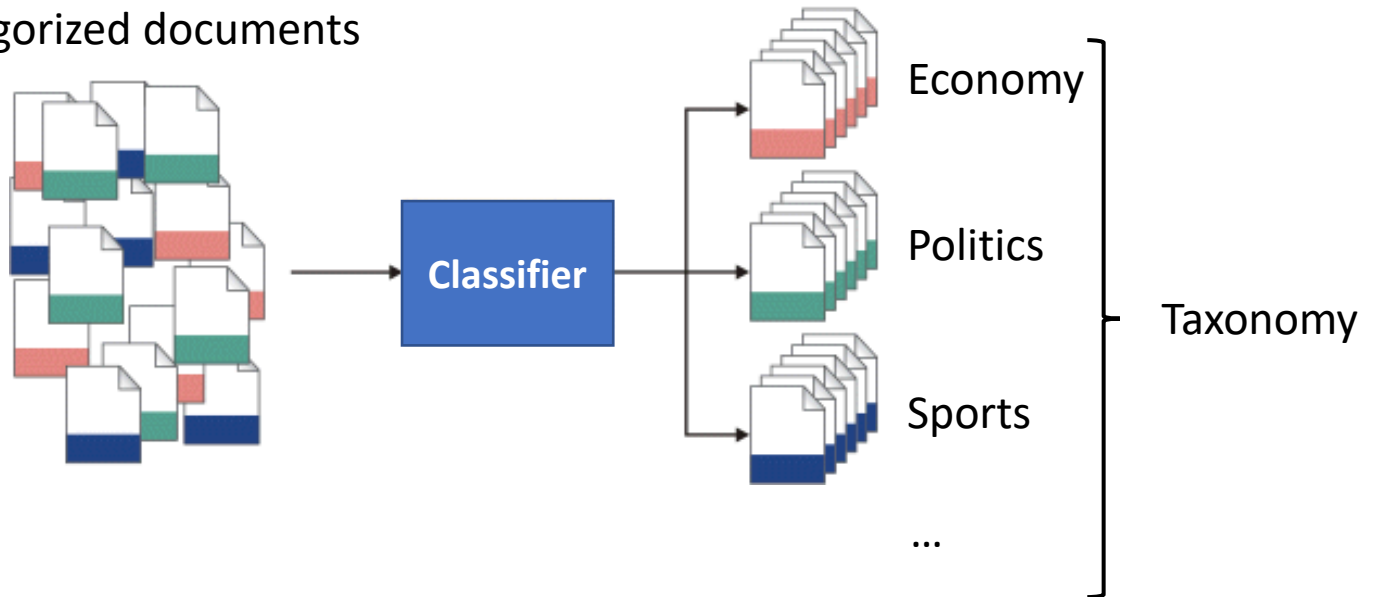
- Wikipedia contains large amounts of information largely unstructured but structured as a taxonomy.
- **DBPedia** aims to create a rigorous database out of Wikipedia.
- A key application is to link data to Wikipedia entries.

<https://en.wikipedia.org/wiki/Portal:Contents>




Document Topic Classification

Uncategorized documents



Classification task


- For new unseen documents, we wish to classify documents with one of the known classes.
- New documents are represented in some feature space and then a machine learning algorithm classifies the new documents.



Perceptron

- All sample vectors $\mathbf{x}^{(i)}$ have their corresponding label $\mathbf{y}^{(i)} = \{+1, -1\}$
- **The perceptron performs a binary prediction \hat{y} based on the observed data x :**


$$\hat{y} = f(x) = \begin{cases} +1 & , \text{if } x_2 \geq m \cdot x_1 + b \\ -1 & , \text{if } x_2 < m \cdot x_1 + b \end{cases}$$



Perceptron

- All sample vectors $\mathbf{x}^{(i)}$ have their corresponding label $\mathbf{y}^{(i)} = \{+1, -1\}$
- **The perceptron performs a binary prediction \hat{y} based on the observed data x :**

$$\hat{y} = f(x) = \begin{cases} +1 & , \text{if } x_2 \geq m \cdot x_1 + b \\ -1 & , \text{if } x_2 < m \cdot x_1 + b \end{cases} = \begin{cases} +1 & , \text{if } 0 \geq m \cdot x_1 + b - x_2 \\ -1 & , \text{if } 0 < m \cdot x_1 + b - x_2 \end{cases}$$




Model error

- The Mean Square Error (MSE) measures the error between the true labels and the predicted labels



$$MSE = \frac{1}{N} \sum_i^N (error_i)^2$$

$$error_i = true_label_i - predictedLabel_i$$



Minimizing the error

$$\text{MeanSquareError} = \frac{1}{\text{TotalSamples}} \sum_i^{\text{TotalSamples}} (\text{label}_i - \text{predictedLabel}_i)^2$$



Learning to minimize the model error

- Initialize the model with random weights
- Compute the model predictions
- Compute the error of each prediction
- Update the model with the samples incorrectly classified.

True label	Predicted label	Error	Update
-1	-1	0	0
-1	+1	-1	$-1 * x$
+1	-1	+1	$+1 * x$
+1	+1	0	0

Learning algorithm

```
[ ]: b=0
      m=0
      model = [m,b]

      max_iters = 30
      mean_square_error = []
      for iter in range(0,max_iters):

          # Compute the model predictions
          predicted_labels = ((observations_x2 - m*observations_x1 - b ) >= 0)*2-1

          # Compute the model error
          error_of_all_samples = (true_labels-predicted_labels)/2

          # Update the model parameters
          update_m = np.mean(error_of_all_samples*observations_x1)
          update_b = np.mean(error_of_all_samples)

          m = m - update_m*0.1
          b = b - update_b*0.1
```

$$\hat{y} = f(x) = \begin{cases} +1 & , \text{if } x_2 - m \cdot x_1 - b \geq 0 \\ -1 & , \text{if } x_2 - m \cdot x_1 - b < 0 \end{cases}$$


$$error = (y - \hat{y})/2 = \begin{cases} +1 \\ 0 \\ -1 \end{cases}$$

$$update_m = error \cdot x_1$$


$$m = m - update_m \cdot learning_{rate}$$

Perceptron learning example


iter = 1



iter = 15



iter = 36




Perceptron: general formulation

- **Binary classification:**

$$z = w_0 + w_1x_1 + \dots + w_nx_n$$


$$\hat{y} = \sigma(z) = \begin{cases} +1 & , \text{if } z \geq 0 \\ -1 & , \text{if } z < 0 \end{cases}$$




- **Input:** Vectors $\mathbf{x}^{(j)}$ and labels $\mathbf{y}^{(j)}$
 - Vectors $\mathbf{x}^{(j)}$ are real valued where $\|\mathbf{x}\|_2 = 1$
- **Goal:** Find vector $\mathbf{w} = (w_1, w_2, \dots, w_d)$
 - Each w_i is a real number

Activation functions

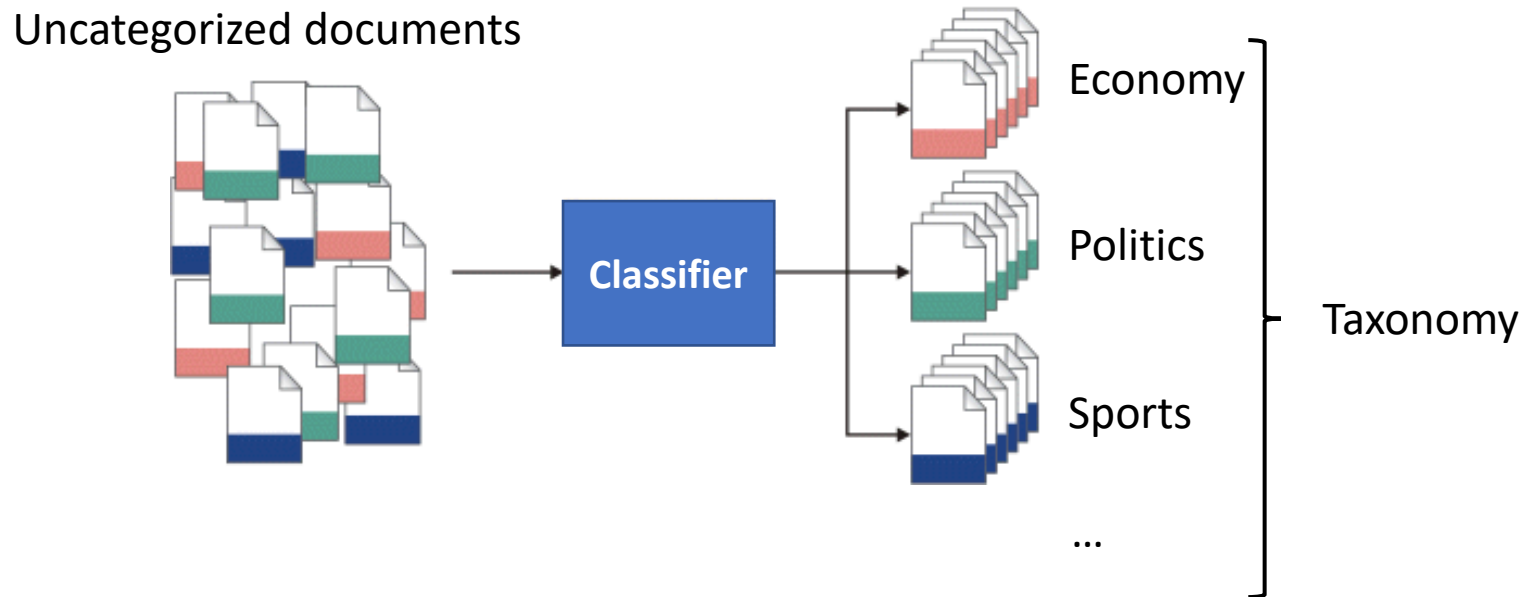
- The perceptron was initially proposed with the step function.
- Historically, other activation functions have been studied.
- It can be shown that the perceptron with the sigmoid activation function corresponds to the logistic regression model.




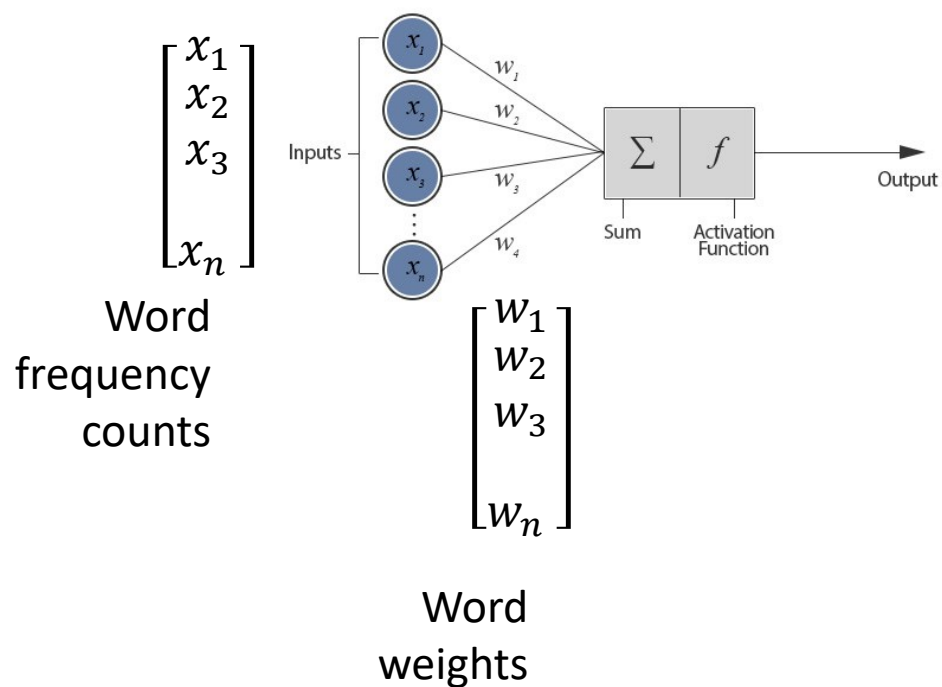
Activation functions



Training Document Categorizers



The sigmoid activation function



Real-world model training

- Robustly training a model for Web data is a complex task.
- In most of the cases, we will use pre-trained models.
- These models were trained on large-scale data.
- These pre-trained models are robust and reliable.

Which and how many are detectable?

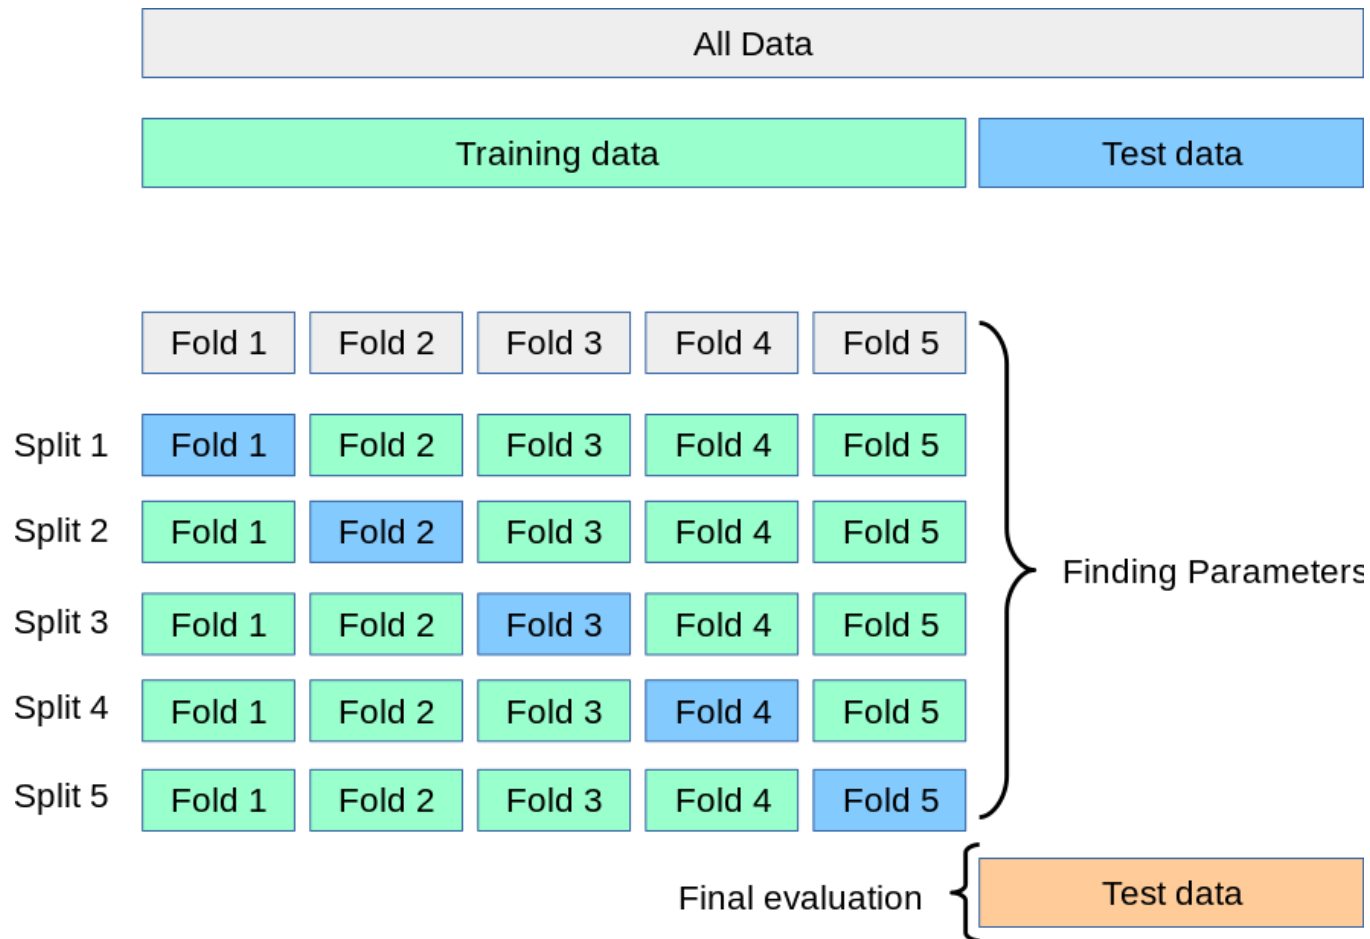
- An important question to ask is which and how many items of the taxonomy are detectable in data?
- A few (well separated ones)? -> Easy!
- A zillion closely related ones? -> Not so easy...
 - Think: Yahoo! Directory, Library of Congress classification, legal applications
 - Quickly gets difficult!
 - Classifier combination is always a useful technique
 - Voting, bagging, or boosting multiple classifiers
 - Much literature on hierarchical classification
 - Definitely helps for scalability, even if not in accuracy
 - May need a hybrid automatic/manual solution



Taxonomies and classification

- In practice, only a few elements of the taxonomy should be used as classes for classification
 - Only the ones offering a stable document class representation.
- The ultimate goal is to link information to an entry on a taxonomy capturing the target domain.
- Ultimately more complete domain representation should be used, e.g. an ontology.

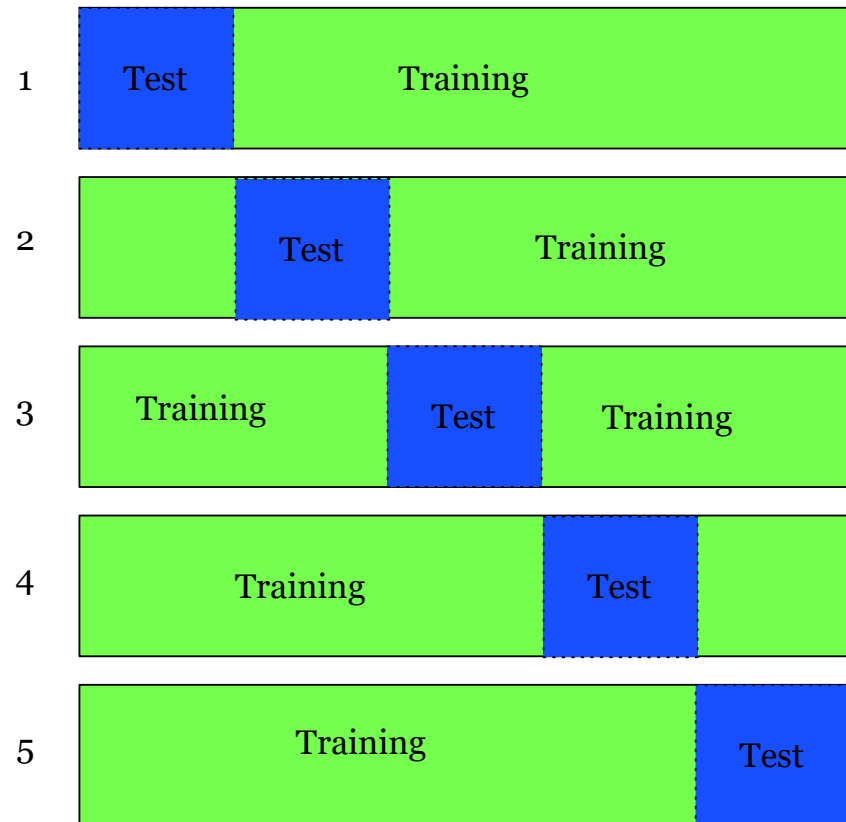
Cross-Validation with held-out test data



Cross-Validation with limited data

- Break up data into 10 folds
 - (Equal positive and negative inside each fold?)
- For each fold
 - Choose the fold as a temporary test set
 - Train on 9 folds, compute performance on the test fold
- Report average performance of the 10 runs

Iteration



Per-class evaluation measures

		Ground-truth	
		True	False
Method	True	True positive	False positive
	False	False negative	True negative

- **Recall:** Fraction of docs in class i classified correctly:

$$Recall = \frac{truePos}{truePos + falseNeg}$$

- **Precision:** Fraction of docs assigned class i that are actually about class i :

$$Precision = \frac{truePos}{truePos + falsePos}$$

- **Accuracy:** Fraction of docs classified correctly:

$$Accuracy = \frac{truePos + trueNeg}{truePos + falsePos + trueNeg + falseNeg}$$

Micro- vs. Macro-Averaging

- If we have more than one class, how do we combine multiple performance measures into one quantity?
- **Macroaveraging:** Compute performance for each class, then average.
- **Microaveraging:** Collect decisions for all classes, compute contingency table, evaluate.

Micro- vs. Macro-Averaging: Example

Class 1

	Truth: yes	Truth: no
Classifier : yes	10	10
Classifier : no	10	970

Class 2

	Truth: yes	Truth: no
Classifier: yes	90	10
Classifier: no	10	890

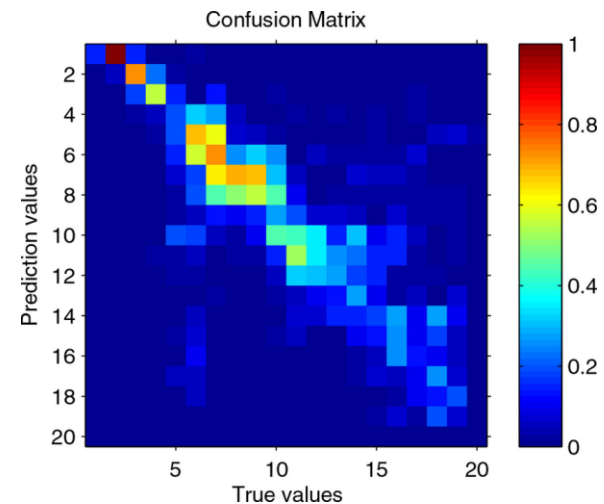

Micro Ave. Table

	Truth: yes	Truth: no
Classifier: yes	100	20
Classifier: no	20	1860

- Macroaveraged precision: $(0.5 + 0.9)/2 = 0.7$
- Microaveraged precision: $100/120 = .83$
- Microaveraged score is dominated by score on common classes

Good practice: Make a confusion matrix

- This (i, j) entry means 53 of the docs actually in class i were put in class j by the classifier.



- In a perfect classification, only the diagonal has non-zero entries
- Look at common confusions and how they might be addressed

Sentiment Classification



THE AMAZING SPIDER-MAN

PG-13 , 2h 16m
[adventure](#) , [mystery and thriller](#) , [action](#)
Directed By: [Marc Webb](#)
In Theaters: [Jul 3, 2012](#) [Wide](#)
Streaming: [Sep 1, 2014](#)
[Marvel Studios](#)



The Amazing Spider-Man: Trailer - Four Minute Super Preview

4 minutes 2 seconds
Added: Feb 8, 2017



The Amazing Spider-Man: Official Clip - Those Are the Best Kind

1 minute 47 seconds
Added: Nov 5, 2019



The Amazing Spider-Man: Official Clip - High School Attack

2 minutes 59 seconds
Added: Nov 5, 2019

[VIEW ALL VIDEOS \(10\)](#)

THE AMAZING SPIDER-MAN REVIEWS

[All Critics](#)

[Top Critics](#)

[All Audience](#)

[NEXT →](#)



Daniel K

★★★★★

Sep 19, 2020

For years I had the blu ray sitting in my collection. I had seen amazing spider man 2 multiple times. I can confirm this movie is amazing. The whole time I was watching I wasn't bored. The movie is entertaining and the score is phenomenal. The first 45 minutes is slow but once uncle Ben dies (spoiler alert) it picks right up. Andrew Garfield is Spider-Man. At first he uses his powers to go after uncle Bens killer but that plot line is abandon and I didn't mind it. After he saves a little boy in a car wreck he takes his...

[Show More](#)



Jacob B

★★★★☆

Sep 14, 2020

As a whole, I like it. Some of the high school scenes are a little too goofy for me and quite honestly don't feel like Marc Webb directed them at all. It also feels as if some important content was missing from the movie (which has been confirmed). Other than that, I like the movie.



France Carl C

★★★★★

Sep 10, 2020

Andrew Garfield is the best Spider-Man



Samantha S

★★★★☆

Sep 04, 2020

Also confused for apparently not having rated this.



E C

★★★★☆

Sep 02, 2020

I thought this film was decent, but it was definitely not as good as the Sam Raimi "Spider-Man" films, and I also didn't think it was as good as the ones with Tom Holland. I thought this film was way too dark in some scenes in my opinion.



Sentiment Classification in Movie Reviews

- Polarity detection:
 - Is an IMDB movie review positive or negative?
- Data: Polarity Data 2.0:
 - <http://www.cs.cornell.edu/people/pabo/movie-review-data>

Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. 2002. Thumbs up? Sentiment Classification using Machine Learning Techniques. EMNLP-2002, 79—86.

Bo Pang and Lillian Lee. 2004. A Sentimental Education: Sentiment Analysis Using Subjectivity Summarization Based on Minimum Cuts. ACL, 271-278

IMDB data in the Pang and Lee database



when _star wars_ came out some twenty years ago , the image of traveling throughout the stars has become a commonplace image . [...]

when han solo goes light speed , the stars change to bright lines , going towards the viewer in lines that converge at an invisible point .

cool .

october sky offers a much simpler image—that of a single white dot , traveling horizontally across the night sky . [. . .]



“ snake eyes ” is the most aggravating kind of movie : the kind that shows so much potential then becomes unbelievably disappointing .

it's not just because this is a brian depalma film , and since he's a great director and one who's films are always greeted with at least some fanfare .

and it's not even because this was a film starring nicolas cage and since he gives a brauvara performance , this film is hardly worth his talents .

Baseline Algorithm (adapted from Pang and Lee)

- Tokenization
- Feature Extraction
- Classification using different classifiers
 - Naïve Bayes
 - MaxEnt
 - SVM

Sentiment Tokenization Issues

- Deal with HTML and XML markup
- Twitter mark-up (names, hash tags)
- Capitalization (preserve for words in all caps)
- Phone numbers, dates, emoticons
- Useful code:
 - [Christopher Potts sentiment tokenizer](#)
 - [Brendan O'Connor twitter tokenizer](#)

REGEX emoticons

[<>]?	# optional hat/brow
[:;=8]	# eyes
[\-o*\']?	# optional nose
[\)\]\(\[dDpP/\:}\{\@\\ \\]	# mouth
	#### reverse orientation
[\)\]\(\[dDpP/\:}\{\@\\ \\]	# mouth
[\-o*\']?	# optional nose
[:;=8]	# eyes
[<>]?	# optional hat/brow

Extracting Features for Sentiment Classification

- How to handle negation

*I **didn't** like this movie*

vs

I really like this movie

- Which words to use?
 - Only adjectives
 - All words
 - All words turns out to work better, at least on this data

Negation

Add NOT_ to every word between negation and following punctuation:

didn't like this movie , but I



didn't NOT_like NOT_this NOT_movie but I

Das, Sanjiv and Mike Chen. 2001. Yahoo! for Amazon: Extracting market sentiment from stock message boards. In Proceedings of the Asia Pacific Finance Association Annual Conference (APFA).
Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. 2002. Thumbs up? Sentiment Classification using Machine Learning Techniques. EMNLP-2002, 79—86.

Word representations

- Binary seems to work better than full word counts
- **Sentiment lexicons** (e.g. SentiWordNet <http://sentiwordnet.isti.cnr.it/>)
 - All WordNet synsets automatically annotated for degrees of positivity, negativity, and neutrality/objectiveness
 - [estimable(J,3)] “may be computed or estimated”
 - Pos 0 Neg 0 Obj 1
 - [estimable(J,1)] “deserving of respect or high regard”
 - Pos .75 Neg 0 Obj .25
- **Affective lexicons**
 - joy–sadness
 - anger–fear
 - trust–disgust
 - anticipation–surprise

Problems:

What makes reviews hard to classify?

- Subtlety:

- Perfume review in *Perfumes: the Guide*:

- “If you are reading this because it is your darling fragrance, please wear it at home exclusively, and tape the windows shut.”

- Dorothy Parker on Katherine Hepburn

- “She runs the gamut of emotions from A to B”

Thwarted Expectations and Ordering Effects

- “This film should be **brilliant**. It sounds like a **great** plot, the actors are **first grade**, and the supporting cast is **good** as well, and Stallone is attempting to deliver a good performance. However, it **can’t hold up**.”
- Well as usual Keanu Reeves is nothing special, but surprisingly, the **very talented** Laurence Fishbourne is **not so good** either, I was surprised.

Summary

- Document topic categorization
- Perceptron and sigmoid function
- Model training
- Sentiment classification